

ARCHITECTURE

INDEXING PIPELINE - OFFLINE SETUP



Domain Documents



Text chunking



Embedding Model



Vector Database

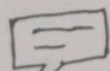
→ Collect PDFs, text files, or web pages

→ Split large documents into smaller chunks.

→ Convert each text chunk into a numerical vector (embedding)

→ Store these vectors in a database like chroma or FAISS for fast searching.

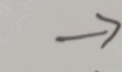
QUERY PIPELINE - EVERY USER QUESTION



User Query



Embedding Model




→ The user enters a question into a vector.

→ Find the most similar document chunks from the vector data base.



Vector Search



→ Retrieve the most relevant chunks


Vector Database → Combine the user's question with the retrieved context.

  → Generate an accurate answer using the context
Prompt Construction

 → Return the final answer to the user
Response.

FLOW CHART

